



Professional forecasters and real-time forecasting with a DSGE model



Frank Smets^{a,b}, Anders Warne^{b,*}, Rafael Wouters^c

^a KU Leuven - Centre for Economic Studies, Leuven, Belgium

^b European Central Bank - Directorate General Research, Kaiserstrasse 29, Frankfurt am Main 60311, Germany

^c National Bank of Belgium, Brussels, Belgium

ARTICLE INFO

Keywords:

Bayesian methods
Real-time data
Survey of Professional Forecasters
Macroeconomic forecasting
Euro area

ABSTRACT

This paper analyses the real-time forecasting performance of the New Keynesian DSGE model of Galí, Smets and Wouters (2012), estimated on euro area data. It investigates the extent to which the inclusion of forecasts of inflation, GDP growth and unemployment by professional forecasters improve the forecasting performance. We consider two approaches for conditioning on such information. Under the “noise” approach, the mean professional forecasts are assumed to be noisy indicators of the rational expectations forecasts implied by the DSGE model. Under the “news” approach, it is assumed that the forecasts reveal the presence of expected future structural shocks in line with those estimated in the past. The forecasts of the DSGE model are compared with those from a Bayesian VAR model, an AR(1) model, a sample mean and a random walk.

© 2014 International Institute of Forecasters. Published by Elsevier B.V. All rights reserved.

1. Introduction

Following the seminal work of Croushore and Stark (2001) on constructing a real-time data set for the US economy, it has become standard to use real-time data when analysing the out-of-sample forecast performances of alternative empirical macromodels.¹ With a few exceptions, much less real-time data analysis has been done on the euro area, partly because a comprehensive real-time euro area data set has only recently become available.² This paper uses the European Central Bank (ECB) real-time data base (RTDB) – described by Giannone, Henry, Lalik, and Modugno (2012) and available on the ECB’s website – to perform two types of analysis.

In this paper we investigate the forecasting performance of the Galí, Smets, and Wouters (2012, GSW) model in real time over the EMU period, and compare it with four alternative non-structural linear models. The GSW model is a version of the model developed by Smets and Wouters (2003, 2007) which has been shown to forecast reasonably well. It is therefore of interest to determine the extent to which these results are robust to the real-time nature of the underlying data in the euro area. A similar exercise was performed recently on US data by Edge and Gürkaynak (2010).

Moreover, we analyse the extent to which professional forecasters’ forecasts of euro area GDP growth, inflation and unemployment (from the ECB’s Survey of Professional Forecasters, SPF) can improve the forecast performance of the DSGE model. We consider two interpretations. Under the “noise” interpretation, the mean professional forecasts are assumed to be noisy indicators of the rational expectations forecasts implied by the DSGE model. Under the “news” interpretation, it is assumed that the forecasts reveal the presence of expected future structural shocks

* Corresponding author. Tel.: +49 69 13448737; fax: +49 69 13446575.
E-mail address: anders.warne@ecb.europa.eu (A. Warne).

¹ See, for example, Croushore (2011) and the literature review on real-time data analysis compiled by Dean Croushore at https://facultystaff.richmond.edu/~dcrousho/docs/realtime_lit.pdf. For an early real-time forecasting exercise, see Diebold and Rudebusch (1991).

² Two exceptions are the studies by Coenen, Levin, and Wieland (2005) and Coenen and Warne (2013).

which are in line with those estimated in the past. This exercise is similar to that performed by Del Negro and Schorfheide (2013) for the United States.

Two sets of results are worth highlighting. First, we find that the point forecasts of the benchmark GSW model can be improved in a mean squared sense via a BVAR, particularly for consumption and real wages, where the GSW model systematically overpredicts real wage growth and underpredicts consumption. These variables are also predicted poorly under the noise and news versions of the GSW model relative to the BVAR model. At the same time, the inflation forecasts from the models using the SPF data are improved considerably relative to the BVAR, making the overall picture more difficult to assess, and likewise when comparing the point forecasts of the DSGE models to common univariate non-structural models.

Second, when comparing the root mean squared errors from the point forecasts of the news version of the GSW model to the benchmark GSW model, also using formal tests, we find that inflation and real wage forecasts are improved over the one- to four-quarter-ahead horizons, while the one- and two-quarter-ahead point forecasts of real GDP growth are also improved when using the 1- and 2-year-ahead SPF data as pure conditioning information. The main forecasting cost is the relative deterioration in the interest rate forecasts. For the noise version, however, it is mainly the inflation forecasts that are improved relative to the benchmark GSW model, while the interest rate forecasts are aggravated. Overall, therefore, we find that the GSW model forecasts for the euro area can be improved by using the SPF data, provided that these data are included based on the news interpretation.

The rest of the paper is structured as follows. Section 2 presents the GSW model. Section 3 presents the real-time data base including the Survey of Professional Forecasts. Section 4 discusses the full-sample estimation results of the benchmark GSW model and provides a brief comparison with the findings for the United States reported by Galí et al. (2012). Section 5 contains the findings of the real-time forecast comparison exercise. Finally, Section 6 summarises the main findings and concludes.

2. The Galí–Smets–Wouters model

This section describes the log-linearized equilibrium conditions of the GSW model. It is a standard medium-sized DSGE model with sticky prices and wages that can explain the main macroeconomic time series, such as output and inflation, and is very similar to the work of Smets and Wouters (2007, SW). One main difference is that it models the labor supply decision on the extensive margin (whether to work or not) rather than the intensive margin (how many hours to work), which allows us to include unemployment as an observable variable.

The model includes eight exogenous shocks: a neutral, factor-augmenting productivity shock ($\hat{\varepsilon}_t^a$), a labor supply shock ($\hat{\varepsilon}_t^s$), a price markup shock ($\hat{\varepsilon}_t^p$), a wage markup shock ($\hat{\varepsilon}_t^w$), a risk premium shock ($\hat{\varepsilon}_t^r$), an exogenous spending shock ($\hat{\varepsilon}_t^g$), an investment-specific technology shock ($\hat{\varepsilon}_t^i$), and a monetary policy shock ($\hat{\varepsilon}_t^m$). In addition, eight observable variables are used to estimate the model. Next, we describe the main structural equations, with E_t denoting the

rational expectations operator conditional on the information at time t , and $\hat{\cdot}$ denoting the deviation of the variable from its steady state growth path.

- Consumption Euler equation. Consumption, \hat{c}_t , depends on lagged consumption because of habit formation and expected future consumption, as well as on the expected short-term real interest rate, $(\hat{r}_t - E_t \hat{\pi}_{t+1} - \hat{\varepsilon}_t^b)$:

$$\hat{c}_t = c_1 E_t [\hat{c}_{t+1}] + (1 - c_1) \hat{c}_{t-1} - c_2 (\hat{r}_t - E_t \hat{\pi}_{t+1} - \hat{\varepsilon}_t^b),$$

where \hat{r}_t is the short-term nominal interest rate and $\hat{\pi}_t$ is the inflation rate, with $c_1 = 1/(1 + (h/\tau))$, $c_2 = (1 - (h/\tau))/(1 + (h/\tau))$, where h is the external habit parameter, τ is the trend growth rate, and $\hat{\varepsilon}_t^b$ is the exogenous AR(1) risk premium process.

- Investment Euler equation. Investment, \hat{i}_t , also depends on past and expected future investment, as well as on the value of capital, \hat{q}_t^k :

$$\hat{i}_t = i_1 \hat{i}_{t-1} + (1 - i_1) E_t \hat{i}_{t+1} + i_2 \hat{q}_t^k + \hat{\varepsilon}_t^q,$$

with $i_1 = 1/(1 + \beta)$, $i_2 = i_1/(\tau^2 \varphi)$, where β is the discount factor, φ is the elasticity of the capital adjustment cost function, and $\hat{\varepsilon}_t^q$ is the exogenous AR(1) process for the investment-specific technology.

- Value of the capital stock. The value of the capital stock is determined by an arbitrage equation which equalizes the expected return on holding capital to the expected real interest rate:

$$\hat{q}_t^k = -(\hat{r}_t - E_t \hat{\pi}_{t+1} - \hat{\varepsilon}_t^b) + q_1 E_t \hat{r}_{t+1}^k + (1 - q_1) E_t \hat{q}_{t+1}^k,$$

with \hat{r}_t^k being the rental rate on capital, $q_1 = r^k/(r^k + (1 - \delta))$, where r^k is the steady-state rental rate on capital and δ is the depreciation rate.

- Goods market equilibrium. In equilibrium, aggregate demand – consisting of consumption, investment, the resources spent on adjusting capital utilization (\hat{v}_t), and an exogenous demand component (spending shock) – has to equal aggregate supply. The latter is determined by a standard Cobb–Douglas production function in effective capital services, k_t , and hours worked, \hat{n}_t :

$$\hat{y}_t = c_y \hat{c}_t + i_y \hat{i}_t + v_y \hat{v}_t + \hat{\varepsilon}_t^g, \\ = \phi_p (\alpha \hat{k}_t + (1 - \alpha) \hat{n}_t + \hat{\varepsilon}_t^a),$$

where \hat{y}_t is output, $c_y = 1 - i_y - g_y$ is the steady-state consumption–output ratio, g_y is the steady-state exogenous spending to output ratio, $i_y = (\tau + \delta - 1)k_y$ is the steady-state investment–output ratio, k_y is the steady-state capital–output ratio, and $v_y = r^k k_y$. The parameter ϕ_p reflects the fixed costs in production, which is assumed to correspond to the price markup in steady state, while $\hat{\varepsilon}_t^g$ and $\hat{\varepsilon}_t^a$ are the AR(1) processes representing exogenous demand components and the TFP process.³

³ The innovation of the TFP process enters the process describing exogenous spending with the parameter ρ_{ga} ; see Table 2 in Section 4.

- Price-setting under the Calvo model with indexation. Inflation is sticky and depends on past and expected future inflation, as well as on the difference between the average ($\hat{\mu}_{p,t}$) and natural ($\hat{\mu}_{p,t}^n$) price markups:

$$\hat{\pi}_t - \gamma_p \hat{\pi}_{t-1} = \pi_1 (E_t \hat{\pi}_{t+1} - \gamma_p \hat{\pi}_t) - \pi_2 (\hat{\mu}_{p,t} - \hat{\mu}_{p,t}^n),$$

with $\pi_1 = \beta$, $\pi_2 = (1 - \theta_p \beta)(1 - \theta_p)/[\theta_p(1 + (\phi_p - 1)\varepsilon_p)]$, with θ_p and γ_p being the probability of price changes and price indexation of the Calvo model, respectively, and ε_p being the curvature of the aggregator function. The average price markup is equal to the inverse of the real marginal cost $\hat{m}c_t = (1 - \alpha)(\hat{w}_t - \hat{p}_t) + \alpha \hat{r}_t^k + \hat{\varepsilon}_t^a$, which is determined by the real wage ($\hat{w}_t - \hat{p}_t$) and the rental rate on capital. The natural price markup is equal to $100\hat{\varepsilon}_t^p$, i.e., it is proportional to the price markup shocks. These shocks are assumed to follow an exogenous ARMA(1, 1) process.

- Wage setting under the Calvo model with indexation. Wage inflation depends on the expected future wage inflation, as well as on past and current inflation, due to partial wage indexation, while nominal wage stickiness implies that it reacts gradually to the difference between the average ($\hat{\mu}_{w,t}$) and natural ($\hat{\mu}_{w,t}^n$) wage markups:

$$\Delta \hat{w}_t = \gamma_w \hat{\pi}_{t-1} + \beta E_t (\Delta \hat{w}_{t+1} - \gamma_w \hat{\pi}_t) - w_1 (\hat{\mu}_{w,t} - \hat{\mu}_{w,t}^n),$$

with Δ being the first difference operator, $w_1 = (1 - \beta\theta_w)(1 - \theta_w)/\theta_w(1 + \varepsilon_w\omega)$, θ_w and γ_w being the probability of wage changes and wage indexation of the Calvo model respectively, ω being the inverse elasticity of labor supply, and ε_w being the curvature of the aggregator function.

- Average and natural wage markups and unemployment. The wage markup is defined as the difference between the real wage and the marginal rate of substitution, which is a function of the smoothed trend of consumption, \hat{z}_t , employment, \hat{e}_t , and the labor supply shock:

$$\hat{\mu}_{w,t} = \hat{w}_t - \hat{p}_t - (\hat{z}_t + \hat{\varepsilon}_t^s + \omega \hat{e}_t),$$

$$= \omega \hat{u}_t.$$

$$\hat{\mu}_{w,t}^n = 100\hat{\varepsilon}_t^w,$$

$$= \omega \hat{u}_t^n.$$

$$\hat{z}_t = (1 - \nu)\hat{z}_{t-1} + \frac{\nu}{1 - (h/\tau)} \left[\hat{c}_t - \frac{h}{\tau} \hat{c}_{t-1} \right],$$

where \hat{u}_t is the unemployment rate, \hat{u}_t^n is the natural rate of unemployment (the unemployment rate that would prevail in the absence of nominal wage rigidities), $\hat{\varepsilon}_t^w$ is assumed to be an exogenous ARMA(1, 1) process, $\hat{\varepsilon}_t^s$ is an AR(1) process representing an exogenous labor supply shock, and ν is a parameter capturing the short-run wealth effects on labor supply. The labor force is given by $\hat{l}_t = \hat{e}_t + \hat{u}_t$.

- Capital accumulation equation. The capital stock, \hat{k}_t , is determined by its lagged value, investment, and the investment-specific technology shock:

$$\hat{k}_t = \kappa_1 \hat{k}_{t-1} + (1 - \kappa_1) \hat{i}_t + \kappa_2 \hat{\varepsilon}_t^q,$$

with $\kappa_1 = (1 - \delta)/\tau$, and $\kappa_2 = (\tau + \delta - 1)(1 + \beta)\tau\varphi$. Capital services used in production is defined as: $\hat{k}_t = \hat{v}_t + \hat{k}_{t-1}$, where \hat{v}_t is capital utilization.

- Optimal capital utilization condition. The degree of capital utilization depends positively on the rental rate on capital:

$$\hat{v}_t = \frac{1 - \psi}{\psi} \hat{r}_t^k,$$

where ψ is the elasticity of the capital utilization cost function.

- Optimal capital/labor input condition:

$$\hat{k}_t = \hat{w}_t - \hat{p}_t - \hat{r}_t^k + \hat{n}_t.$$

- Monetary policy rule:

$$\hat{r}_t = \rho_R \hat{r}_{t-1} + (1 - \rho_R) (r_\pi \hat{\pi}_t + r_y \hat{y}_t^{\text{gap}} + r_{\Delta y} \Delta \hat{y}_t^{\text{gap}}) + \hat{\varepsilon}_t^r,$$

with $\hat{y}_t^{\text{gap}} = \hat{y}_t - \hat{y}_t^{\text{flex}}$, the difference between actual output and the output in the flexible price and wage economy, i.e., in the absence of distorting price and wage markup shocks.

As productivity is written in terms of hours worked, we also introduce an auxiliary equation to link the observed total employment (\hat{e}_t) to unobserved hours worked, as per SW (2003):

$$\hat{e}_t - \hat{e}_{t-1} = E_t \hat{e}_{t+1} - \hat{e}_t + \frac{(1 - \beta\theta_e)(1 - \theta_e)}{\theta_e} (\hat{n}_t - \hat{e}_t),$$

where θ_e is the fraction of firms that are able to adjust employment to its desired total labor input.

The model is consistent with a balanced steady-state growth path, driven by deterministic labor augmenting trend growth. The observed variables for the euro area are given by quarterly data on the log of real GDP (y_t), the log of real private consumption (c_t), the log of real total investment (i_t), the log of the GDP deflator ($p_{y,t}$), the log of real wages ($w_t - p_{y,t}$), the log of total employment (e_t), the unemployment rate (u_t), and the short-term nominal interest rate (r_t). With all variables except the unemployment rate and the interest rate being measured in first differences, the measurement equations for the euro area are given by:

$$\begin{bmatrix} \Delta y_t \\ \Delta c_t \\ \Delta i_t \\ \pi_{y,t} \\ \Delta w_t - \pi_{y,t} \\ \Delta e_t \\ u_t \\ r_t \end{bmatrix} = \begin{bmatrix} \bar{\tau} + \bar{e} \\ \bar{\tau} + \bar{e} \\ \bar{\tau} + \bar{e} \\ \bar{\pi} \\ \bar{\tau} \\ \bar{e} \\ \bar{u} \\ 4\bar{r} \end{bmatrix} + \begin{bmatrix} \Delta \hat{y}_t \\ \Delta \hat{c}_t \\ \Delta \hat{i}_t \\ \hat{\pi}_t \\ \Delta \hat{w}_t - \Delta \hat{\pi}_t \\ \Delta \hat{e}_t \\ \hat{u}_t \\ 4\hat{r}_t \end{bmatrix}, \quad (1)$$

where $\hat{u}_t = \hat{l}_t - \hat{e}_t$. The steady-state parameters are determined as

$$\bar{\tau} = 100(\tau - 1), \quad \bar{\pi} = 100(\pi - 1),$$

$$\bar{r} = 100 \left(\frac{\pi\tau}{\beta} - 1 \right), \quad \bar{u} = 100 \left(\frac{\phi_w - 1}{\omega} \right),$$

where $(\phi_w - 1)$ is the steady-state labor market markup, π is steady-state inflation, and \bar{e} reflects steady-state labor force growth and is added to the real variables that are not measured in per capita terms.

The following parameters are not identified by the estimation procedure and are therefore calibrated: $g_y = 0.18$, $\delta = 0.025$, and $\varepsilon_p = 10$.

3. The euro area RTDB and the SPF

Following GSW, we estimate the DSGE model using eight macroeconomic time series for the euro area: real GDP, consumption, investment, employment, unit labor costs, GDP deflator inflation, the Euribor rate and the unemployment rate, with the first five being log differenced. Real-time vintages of these data are available for downloading from the ECB's Statistical Data Warehouse and are described by Giannone et al. (2012).⁴ The frequency of the vintages is monthly, corresponding to their publication in the ECB's Monthly Bulletin, and the first vintage starts in January 2001. The latest available vintage we use in this paper is March 2011.

Table 1 presents the time flow of data releases available for the euro area Real-Time Data Base (RTDB) and the Survey of Professional Forecasters (SPF).⁵ We take the vintage of the last month of the quarter, in order to convert the monthly vintages into quarterly vintages. As is clear from the table, this implies that monthly unemployment and HICP inflation are available for the first month of the quarter, whereas the monthly interest rate is available for the first and second months of the quarter. As we need the full quarter of monthly observations to construct the quarterly observation, we ignore the partial information available during the quarter. This implies that quarterly unemployment, HICP inflation and the interest rate are observed with a one-quarter lag. Using the vintage of the last month in the quarter implies that the quarterly series are also typically available with one lag, with the exception of employment and wage compensation, which are only available with a two-quarter lag. In the forecasting exercises of Section 5, we will use the method of Waggoner and Zha (1999) to “nowcast” employment and wages based on information from the same quarter on real GDP and the other variables.⁶

Each monthly data vintage from the RTDB typically only covers data starting in the mid-1990s. To extend the real-time data backwards, we make use of updates of the quarterly database constructed for estimating the Area-Wide Model (AWM). Since 2000, the AWM database has been updated annually; see Fagan, Henry, and Mestre (2005).

Fig. 1 plots the first release and the first annual revision of real GDP growth, GDP deflator inflation and the unemployment rate (left panel), as well as the difference between the first release and the first annual revision (right panel). The standard deviation of the annual revision in real GDP growth lies between 0.1 and 0.2, and is quite persistent. The downward revision was particularly large in the

Table 1

Time flow of data releases available for the RTDB and the SPF over a quarter.

	Quarter			
	Month 1		Month 2	Month 3
	RTDB M1	SPF	RTDB M2	RTDB M3
Monthly series	u_{m-2}	u_{m-2}	u_{m-2}	u_{m-2}
	π_{m-2}	π_{m-1}	π_{m-2}	π_{m-2}
	r_{m-1}	r_{m-1}	r_{m-1}	r_{m-1}
Quarterly series	y_{q-2}	y_{q-2}	y_{q-2}	y_{q-1}
	c_{q-2}	c_{q-2}	c_{q-2}	c_{q-1}
	i_{q-2}	i_{q-2}	i_{q-2}	i_{q-1}
	$p_{y,q-2}$	$p_{y,q-2}$	$p_{y,q-2}$	$p_{y,q-1}$
	e_{q-2}	e_{q-2}	e_{q-2}	e_{q-2}
	w_{q-2}	w_{q-2}	w_{q-2}	w_{q-2}
	u_{q-2}	u_{q-2}	u_{q-1}	u_{q-1}
	r_{q-1}	r_{q-1}	r_{q-1}	r_{q-1}

Note: Unemployment is denoted by u , HICP by π , the average quarterly 3-month nominal interest rate by r , real GDP by y , real private consumption by c , the GDP deflator by p_y , total employment by e , and wages by w .

most recent recession. The variability of the annual revision in inflation is of the same size, but is much less persistent. Finally, revisions in unemployment are the most persistent.

One source of revisions in the euro area data set is the increasing number of EU countries which are members of the euro area. Over the estimation sample, the euro area developed from 12 to 16 members: updates 4, 5, and 6 of the AWM database cover the euro area 12 data and are taken from 2003, 2004, and 2006, respectively. The euro area 13 composition is available in update 7 from 2007, while the euro area 15 composition is available in update 8, dated 2008. The last two updates that we make use of, 9 and 10, both cover the euro area 16 composition, and were frozen in 2009 and 2010. The available files prior to update 7 are dated in September, though the time they were frozen is unknown; as of update 7, the AWM data are frozen at the beginning of August.

Table 1 also shows that the SPF forecasts for HICP inflation, real GDP growth and unemployment typically become available in the first month of the quarter.⁷ We associate this forecast with the quarter. The SPF data set contains average one- and two-year-ahead forecasts covering the period 1999Q1–2010Q4. Due to the different frequencies and lags in the release of HICP inflation, real GDP and unemployment, the end dates of the one- and two-year-ahead forecasts differ across variables. For HICP inflation, the Q1-release one-year-ahead forecasts refer to the annual inflation in December in the same year, the Q2-release refers to March in the following year, etc. For real GDP growth, the “one-year-ahead forecast” in the Q1-release refers to annual growth in the third quarter of the same year, etc. Finally, for the unemployment

⁴ See also the detailed information about the RTDB provided by Giannone, Henry, Lalik, and Modugno (2010).

⁵ See for example Bowles et al. (2007) and Garcia (2003) for information on the ECB's SPF. For a recent study using SPF data, see Genre, Kenny, Meyler, and Timmermann (2013).

⁶ Relative to the vintage date, employment and wages are actually backcasted, while the remaining variables are nowcasted.

⁷ The inflation forecasts in the SPF only cover HICP inflation, not the GDP deflator. We therefore use the HICP inflation forecasts. In the estimation under the noise interpretation, the difference is picked up by the measurement error term. The model under the news interpretation is estimated from the RTDB data only, and SPF forecasts are only used as conditioning information when forecasting.

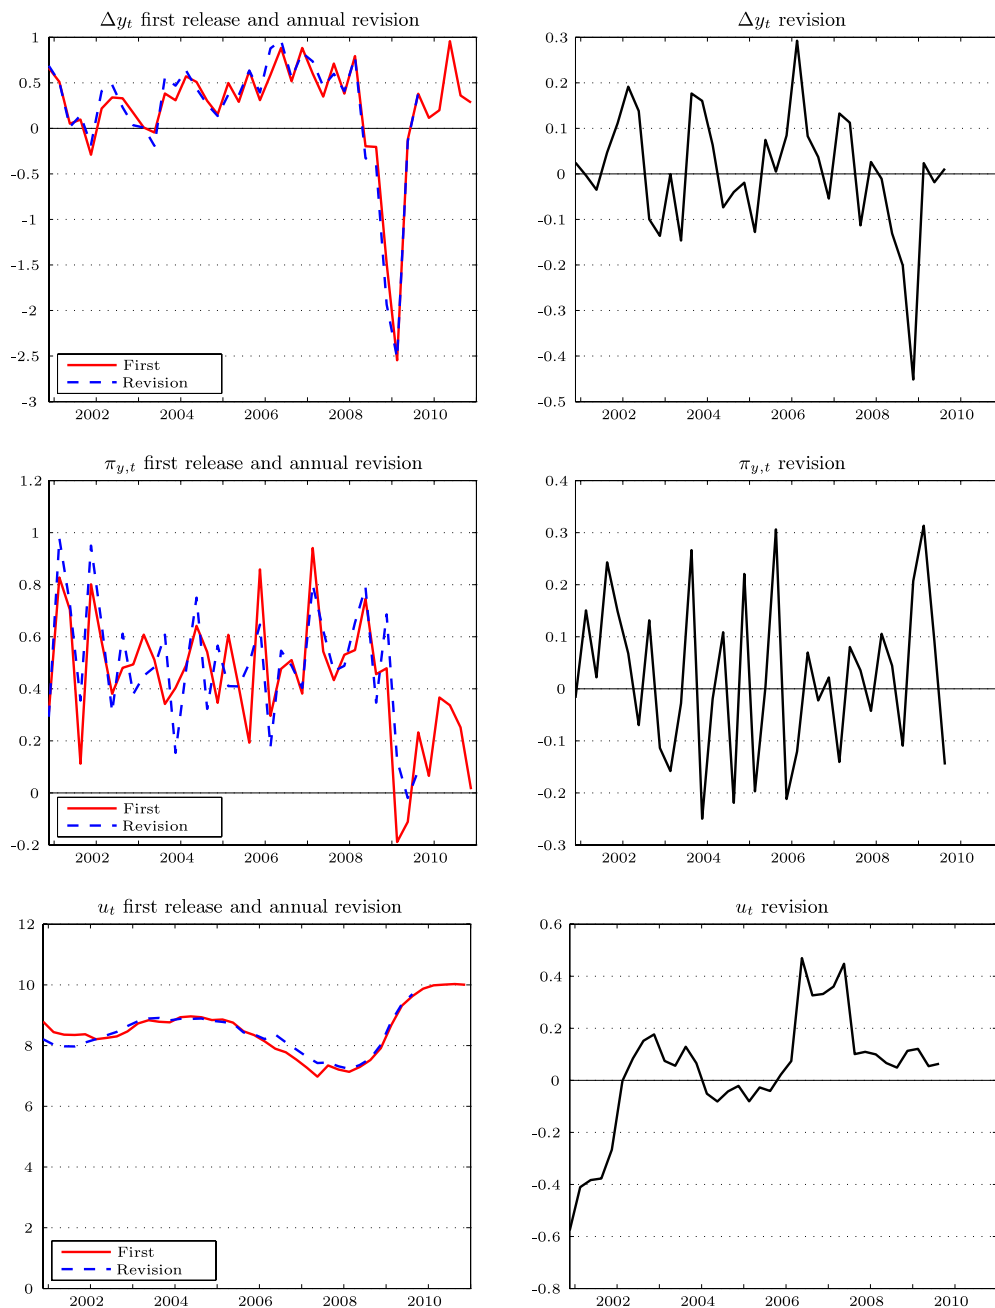


Fig. 1. First release and annual revision data for real GDP growth (Δy_t), GDP deflator inflation ($\pi_{y,t}$), and the unemployment rate (u_t), 2000Q4–2010Q4.

rate, the “one-year-ahead” in the Q1-release refers to the unemployment rate in November of the same year, the Q2-release to the rate in February of the next year, etc. If we take the release quarters as the current date for these forecasts, then we may think of this as having three- and seven-quarter-ahead forecasts for HICP inflation and unemployment and two- and six-quarter-ahead forecasts for real GDP growth.

The information set available to the professional forecasters is smaller than the RTDB available in the last month of the quarter, as the previous quarter’s national account data are not available early in the quarter. On the other

hand, it is clear that the professional forecasters have a lot more information available for nowcasting the last quarter than the data we use from the RTDB. As a result, it is not clear whether the net information advantage is positive or negative.

4. Full-sample estimation results

In this section we first discuss the estimation results using the 2010Q1-vintage data set and make some comparisons with those reported for the United States by GSW

Table 2

Prior distributions and posterior estimates for the US and euro area models.

Parameter	Prior			Posterior							
	Type	Mean	St.dev	United States (1966:1–2007:4)				Euro area (1985:1–2009:4)			
				Mode	Mean	5%	95%	Mode	Mean	5%	95%
Structural parameters											
φ	N	4.0	1.0	4.09	3.96	2.34	5.58	4.65	4.79	3.34	6.31
h	B	0.7	0.1	0.78	0.75	0.65	0.85	0.65	0.64	0.54	0.72
ω	N	2.0	1.0	3.99	4.35	3.37	5.32	5.66	5.56	4.49	6.63
ν	B	0.5	0.2	0.02	0.02	0.01	0.04	0.06	0.12	0.03	0.34
θ_p	B	0.5	0.15	0.58	0.62	0.53	0.71	0.85	0.85	0.79	0.90
θ_w	B	0.5	0.15	0.47	0.55	0.44	0.66	0.74	0.72	0.60	0.89
γ_p	B	0.5	0.15	0.26	0.49	0.20	0.78	0.22	0.27	0.11	0.49
γ_w	B	0.5	0.15	0.16	0.18	0.07	0.29	0.22	0.25	0.12	0.42
ψ	B	0.5	0.15	0.57	0.56	0.36	0.75	0.46	0.48	0.29	0.69
ϕ_p	N	1.25	0.12	1.74	1.74	1.61	1.88	1.48	1.48	1.31	1.65
ϕ_w	N	1.25	0.12	1.18	1.22	1.15	1.29	1.53	1.51	1.41	1.62
α	N	0.3	0.05	0.17	0.17	0.14	0.20	0.22	0.22	0.19	0.26
θ_e	B	0.5	0.15	–	–	–	–	0.71	0.71	0.65	0.76
ρ_R	B	0.75	0.1	0.85	0.86	0.82	0.89	0.86	0.86	0.81	0.89
r_π	N	1.5	0.25	1.91	1.89	1.62	2.16	1.25	1.27	1.02	1.57
r_y	N	0.12	0.05	0.15	0.16	0.11	0.22	0.19	0.19	0.14	0.25
$r_{\Delta y}$	N	0.12	0.05	0.24	0.25	0.20	0.30	0.02	0.02	–0.00	0.06
$\tilde{\pi}$	G	0.62	0.1	0.62	0.66	0.49	0.83	0.55	0.56	0.44	0.70
$\tilde{\beta}$	G	0.25	0.1	0.31	0.31	0.17	0.43	0.24	0.27	0.13	0.43
\bar{l}	N	0.0	2.0	–1.65	–1.52	–3.83	0.77	–	–	–	–
\bar{e}	N	0.2	0.5	–	–	–	–	0.22	0.22	0.20	0.25
$\bar{\tau}$	N	0.4	0.1	0.34	0.34	0.30	0.37	0.14	0.14	0.08	0.20
τ_{wE}	N	0.2	0.1	0.07	0.08	0.03	0.12	–	–	–	–
St.dev. of the innovations											
σ_a	U	2.5	1.44	0.41	0.42	0.37	0.46	0.58	0.60	0.46	0.78
σ_b	U	2.5	1.44	1.73	1.60	0.56	2.50	0.24	0.28	0.16	0.44
σ_g	U	2.5	1.44	0.47	0.48	0.43	0.52	0.30	0.31	0.28	0.35
σ_q	U	2.5	1.44	0.42	0.42	0.34	0.49	0.49	0.49	0.39	0.60
σ_r	U	2.5	1.44	0.21	0.22	0.19	0.24	0.11	0.11	0.10	0.13
σ_p	U	2.5	1.44	0.05	0.11	0.03	0.18	0.35	0.49	0.21	1.02
σ_w	U	2.5	1.44	0.04	0.06	0.01	0.13	0.30	0.76	0.16	3.66
σ_s	U	2.5	1.44	1.07	1.17	0.89	1.45	1.02	1.07	0.85	1.33
Persistence of the exogenous processes: $\rho = \text{AR}(1)$, $\mu = \text{MA}(1)$											
ρ_a	B	0.5	0.2	0.98	0.98	0.97	0.99	0.98	0.98	0.97	0.99
ρ_b	B	0.5	0.2	0.36	0.42	0.19	0.67	0.91	0.91	0.84	0.96
ρ_g	B	0.5	0.2	0.97	0.97	0.96	0.99	0.99	0.99	0.98	1.00
ρ_{ga}	N	0.5	0.25	0.69	0.69	0.55	0.83	0.18	0.19	0.09	0.30
ρ_q	B	0.5	0.2	0.72	0.75	0.62	0.88	0.36	0.35	0.18	0.53
ρ_r	B	0.5	0.2	0.09	0.10	0.02	0.17	0.30	0.30	0.16	0.44
ρ_p	B	0.5	0.2	0.76	0.43	0.07	0.79	0.56	0.53	0.27	0.76
μ_p	B	0.5	0.2	0.59	0.57	0.24	0.96	0.44	0.47	0.25	0.71
ρ_w	B	0.5	0.2	0.99	0.98	0.97	1.00	0.91	0.89	0.81	0.95
μ_w	B	0.5	0.2	0.67	0.63	0.35	0.91	0.85	0.80	0.65	0.90

Note: The prior distribution types are normal (N), standardized beta (B), gamma (G), and uniform (U). The parameter $\tilde{\beta} = 100(\beta^{-1} - 1)$. The parameter ϕ_w has prior mean 1.5 and standard deviation 0.25 for the euro area, while the parameter $\bar{\tau}$ has prior mean 0.3 and standard deviation 0.1 for the vintages prior to 2008 and standard deviation 0.05 thereafter. The US results are taken from Galí et al. (2012).

The uniform priors all have lower bound 0 and upper bound 5. The parameter ρ_{ga} measures the effect of TFP innovations on exogenous spending. The persistence parameter for the labor supply process $\hat{\varepsilon}_t^s$ is calibrated and given by $\rho_s = 0.999$.

(2012). We estimate the model over the period 1985Q1–2009Q4 using Bayesian full-system estimation techniques, as per SW (2003, 2007). The period from 1980Q1 to 1984Q4 is used as the training period.⁸

⁸ Provided that the log-linearized GSW model has a unique and convergent solution for a given value of the parameters, it can be written in state space form. The Kalman filtering and smoothing algorithms that take missing observations into account can then be used for estimation and (conditional) forecasting with state space models; see for example Durbin and Koopman (2012). The same algorithms are also used when

Table 2 reports the parameter estimates and prior distributions that we have used.⁹ A few striking differences from the US results are worth mentioning. First, the av-

the GSW model is extended to take the SPF into account via either the news or noise interpretation.

⁹ Most of the priors we have used are standard in the literature, and, overall, quite uninformative. When there are many parameters to estimate, there are not many good reasons for picking different priors for the euro area and the US. For those where one may argue that the parameters should be different, such as the stickiness parameters, we have opted to use quite uninformative priors.

Table 3

Variance decompositions for the US and euro area models (in percent).

Variance decomposition	Output	Inflation	Employment	Unemployment
<i>10-quarter horizon</i>				
Demand shocks				
Risk premium	6 / 32	2 / 12	16 / 67	20 / 64
Exogenous spending	3 / 0	1 / 0	7 / 1	8 / 0
Investment specific	9 / 2	3 / 0	12 / 2	10 / 1
Monetary policy	5 / 6	8 / 0	11 / 11	11 / 11
Supply shocks				
Productivity	59 / 54	6 / 8	5 / 1	4 / 2
Price markup	2 / 0	27 / 61	3 / 0	0 / 0
Labor market shocks				
Wage markup	6 / 0	53 / 17	18 / 2	41 / 15
Labor supply	11 / 3	0 / 0	29 / 12	5 / 4
<i>40-quarter horizon</i>				
Demand shocks				
Risk premium	2 / 14	1 / 12	6 / 43	7 / 54
Exogenous spending	1 / 0	1 / 0	3 / 4	3 / 0
Investment specific	5 / 1	2 / 0	4 / 1	3 / 1
Monetary policy	2 / 2	5 / 0	4 / 7	4 / 9
Supply shocks				
Productivity	56 / 75	4 / 12	3 / 0	1 / 0
Price markup	1 / 0	18 / 53	1 / 2	0 / 0
Labor market shocks				
Wage markup	17 / 0	67 / 19	39 / 4	80 / 27
Labor supply	17 / 5	0 / 0	40 / 0	2 / 3

Note: The first entry gives the variance decompositions for the US (1966:1–2007:4) from GSW (2012); the second entry gives the decompositions for the euro area (1985:1–2009:4).

average unemployment rate over the period 1985–2009 is quite a bit higher in the euro area (about 9%) than in the United States (5%). In steady state, the unemployment rate is proportional to the wage markup and the labor supply elasticity. For the euro area, the wage markup is estimated to be quite a bit higher¹⁰ (around 50%), and the labor supply elasticity somewhat lower. In other words, the labor supply responds less to changes in real wages in the euro area.

Second, the parameter governing the short-run wealth effects on labor supply, ν , is quite small, as in the United States. Roughly speaking, this amounts to a preference specification closer to that of Greenwood, Hercowitz, and Huffman (1988), in which the wealth effects are close to zero in the short run. As was discussed at length by GSW, this helps to ensure that not only employment, but also the labor force, moves procyclically in response to most shocks.

Third, turning to some of the other parameters that enter the price and wage Phillips curve, the euro area economy appears to be much stickier than the US economy. The estimated degree of price and wage indexation is relatively small in both areas (around 0.25), but the estimated Calvo probabilities of unchanged wages and prices are quite a bit higher. The average wage contract duration is slightly over three quarters, whereas the average duration of unchanged prices is more than six quarters. This is consistent with some of the micro evidence on price and wage adjustment.¹¹

¹⁰ This may reflect a higher degree of unionization and collective wage bargaining in the euro area than in the US; see WDN (2009).

¹¹ See, for instance, Altissimo, Ehrmann, and Smets (2006) and WDN (2009).

Fourth, it is worth pointing out that the monetary policy reaction coefficient for the output gap (defined as the deviation relative to the constant markup output) is quite high (0.19) compared to the United States, whereas the coefficient on inflation is quite a bit lower (though higher than one).

Finally, focusing on the volatility and persistence of the eight structural shocks, the most striking difference is that the risk premium shock is much more persistent in the euro area, whereas the investment-specific technology shock is much less persistent.

Overall, the estimation results for the euro area point to a less flexible economy with more persistence in the effects of various shocks on economic activity, prices and unemployment.

Before turning to the real-time forecasting results, it is also worth discussing briefly the forecast error variance decomposition at the 10- and 40-quarter horizons (Table 3). At the business cycle frequency, about half of the fluctuations in output are driven by demand shocks, and the risk premium shock in particular. The risk premium shock explains almost two thirds of the movement in unemployment at the 2.5-year horizon. The monetary policy shock explains another 11%. The most important shock driving output is the productivity shock. The price inflation is mostly driven by the price markup shock (61%) and the wage markup shock (17%).

In the longer run (after ten years), the role of wage markup shocks in driving both unemployment and inflation becomes more important. However, these shocks are still much less important than in the United States, where they account for 60%–80% of the movements. The role of

demand shocks in explaining real output and unemployment drops somewhat in the longer run, but remains much more important than in the US. Productivity shocks become relatively more important. In the longer run, inflation is mostly driven by price and wage markup shocks.

These full-sample estimation results are very similar when we re-estimate the model using the SPF forecasts as noisy indicators of the model-consistent expectations (see Section 5). We find that the estimates of the standard deviation of the i.i.d. normal measurement error are relatively large: 0.76 for expected annual real GDP growth, 0.32 for expected GDP deflator inflation and 0.60 for the expected unemployment rate.

5. Real-time forecasting performance

In this section we evaluate the real-time forecasting performance of the GSW model over the EMU period and compare it with the performances of seven alternative models. With the exception of three simple non-structural models, each of these models is re-estimated on an annual basis from the first RTDB vintage in 2001Q1 onwards; i.e., the second based on the 2002Q1 vintage, and so on. The forecasts are conditional on the data observed in the last historical period, where the information available in that period is used to backcast the variables that are missing in that period (typically employment and wage compensation). For example, the RTDB vintage 2001Q1 forecasts are computed for 2000Q4–2001Q4, with conditioning assumptions for 2000Q4 based on the historical data available for that quarter.

One question in real-time forecast evaluation exercises is which actual data to use for evaluating the forecasts and calculating the forecast errors. As is common in the literature, we use the first annual revision of the data (as in Fig. 1). We have checked the robustness of our findings against two possible alternatives for the actual data: (1) the first release data and (2) latest vintage data. Overall, the results are very similar.

We compare the point forecasts of the GSW model with those from seven alternative models. The four competing non-structural models are the random walk, the sample mean, an AR(1) model for deviations around the sample mean, and a BVAR model using the same eight observed variables as in the GSW model. The sample mean model is re-estimated in each quarter and uses data from the last 40 quarters. The AR(1) model takes the deviations around the sample mean for the previous model and adds an autoregressive lag to this, and is therefore similar to the AR(1)-gap model used for inflation forecasting by Faust and Wright (2013). The autoregressive parameter for each variable is re-estimated for each RTDB vintage using the data available for that variable.¹²

¹² We have also tested replacing the sample mean for annual real GDP growth, annual GDP deflator inflation, and the unemployment rate with the “five-year-ahead” mean point forecasts of annual real GDP growth, HICP inflation, and the unemployment rate from the SPF vintages. Such data are available for each vintage from 2001Q1 onwards, and for Q1 in 1999 and 2000. The use of the SPF for replacing means was suggested to us by one of the referees. However, the MSEs from these variants

The BVAR estimation follows Villani (2009). It is estimated using a prior on the steady-state mean and standard deviation of the variables, which is the same as the prior steady-state mean and standard deviation used in estimating the DSGE model (with the exception of the standard deviation of unemployment). In addition, a fairly standard Minnesota-type prior with a diffuse prior on the covariance matrix is used.

The benchmark GSW model is also compared with three alternative estimated GSW models in which the mean forecasts of real GDP growth, HICP inflation, and unemployment from the SPF are used as additional information. We consider two possible interpretations of those professional forecasts.¹³ Under the “noise” interpretation, the mean professional forecasts are assumed to be noisy indicators of the rational expectations forecasts implied by the DSGE model. Specifically, mean SPF forecasts of annual inflation are added to the set of measurement equations in Eq. (1):

$$\pi_{t+3|t}^a = 4\bar{\pi} + E_t[\hat{\pi}_{t+3} + \hat{\pi}_{t+2} + \hat{\pi}_{t+1} + \hat{\pi}_t] + \eta_{\pi,t},$$

where $\pi_{t+3|t}^a$ is the mean SPF forecast of annual inflation between $t+3$ and $t-1$ in period t , and $E_t[\hat{\pi}_{t+i}]$ is the rationally expected quarterly inflation rate (in deviation from the steady state) in period $t+i$, using information available until period t in the GSW model. The sum of the first and second terms (in brackets) on the right hand side is therefore equal to the rational expectations forecast of annual inflation three quarters ahead, while $\eta_{\pi,t}$ is an i.i.d. normal measurement error with mean zero and standard deviation σ_{π} . Similarly, measurement equations are added to Eq. (1) for the mean SPF forecast of unemployment three quarters ahead and for the mean SPF forecast of annual real GDP growth two quarters ahead, both with individual measurement errors. The additional randomness makes it possible to estimate the parameters of the GSW model extended with the SPF data.¹⁴ As is discussed in Section 4, the standard deviations of the errors in the measurement equations are quite large. Conditional forecasts are computed using the Waggoner and Zha (1999) approach with hard conditions; see Warne (2013) for details of its implementation in linear state space models.

Under the “news” interpretation, it is assumed that the forecasts reveal the presence of expected future structural

of the sample mean and AR(1)-gap models worsened the forecasting performances of these two models for both real GDP growth and GDP deflator inflation, while the forecasts of the unemployment rate improved somewhat. We therefore decided against using the “five-year-ahead” mean point forecasts from the SPF. Note also that the quotation marks reflect the fact that the five-year-ahead forecasts relate to calendar years, and are therefore not strictly five years ahead for most vintages.

¹³ While the SPF is conducted quarterly and the data we have collected from the RTDB are also quarterly, it would be possible to augment the GSW model with monthly conjunctural data using a mixed frequency approach, as was suggested by Giannone, Monti, and Reichlin (2010). An important aspect of their methodology is that the extra information provided by the monthly conjunctural data is valuable (relative to the augmented model) only because it is more timely. For an application to short-term forecasting of Austrian real GDP, see for example Červená and Schneider (2014).

¹⁴ For inflation expectations, our approach is similar to that of Del Negro and Eusepi (2011), except that we add measurement noise instead of modifying the policy rule.

Table 4

Modified Diebold–Mariano tests of GSW model based cases with SPF conditioning information versus the benchmark GSW model (RTDB vintages 2001Q1–2010Q4).

Variable	Noise				News 1-year				News 1- and 2-year			
	1	2	3	4	1	2	3	4	1	2	3	4
Output	0.22	0.21	0.42	0.48	0.01	0.03	0.20	0.39	0.01	0.04	0.20	0.41
Consumption	0.07	0.24	0.42	0.46	0.02	0.33	0.56	0.66	0.01	0.31	0.58	0.71
Investment	0.62	0.41	0.41	0.44	0.48	0.10	0.04	0.08	0.63	0.23	0.07	0.07
Inflation	0.01	0.01	0.02	0.03	0.03	0.00	0.02	0.05	0.01	0.00	0.02	0.05
Employment	0.16	0.19	0.45	0.49	0.91	0.77	0.36	0.54	0.92	0.80	0.37	0.52
Real wages	0.99	1.00	0.99	1.00	0.03	0.00	0.00	0.00	0.02	0.00	0.00	0.00
Unemployment	0.25	0.20	0.35	0.44	0.37	0.42	0.54	0.63	0.27	0.36	0.51	0.63
Interest rate	0.95	0.89	0.88	0.88	0.98	0.98	0.94	0.88	0.99	0.98	0.96	0.91

Notes: The modified Diebold–Mariano test has been calculated as per Harvey et al. (1997, equation (9)) for the squared forecast errors of a GSW model based case where the SPF is used as conditioning information (noise model with 1-year-ahead SPF; news model with 1-year-ahead SPF; news model with 1- and 2-year-ahead SPF) versus the squared forecast errors of the benchmark GSW model. Percentile values taken from the Student's t -distribution with $N_h - 1$ degrees of freedom are shown above, with N_h being the number of h -step-ahead forecast errors, $N_h = 36 - h$. Smaller percentile values favor models that include the SPF as conditioning data, and larger percentile values favor the DSGE model without this data. Numbers in bold refer to percentile values which are less than or equal to 5%, and numbers in italics to values greater than or equal to 95%.

shocks which are in line with those estimated over the past. This exercise is similar to that performed by Del Negro and Schorfheide (2013) for the United States. In this case, the corresponding DSGE model forecast of the annual real GDP growth two quarters ahead, the annual GDP deflator inflation three quarters ahead, and the unemployment rate three quarters ahead, will be identical to the mean SPF forecast. Again, the Waggoner and Zha methodology is used to compute the conditional forecasts, and we report results for two cases: one in which we only use the one-year-ahead forecasts and another one in which we also use the two-year-ahead SPF forecasts.¹⁵

The forecasting performance exercise below addresses two main questions. First, are the benchmark GSW model forecasts improved upon by utilizing the SPF data? Second, can the GSW models with and without conditioning on the SPF data compete with the non-structural models?

Fig. 2 displays the RMSEs for the three cases of the GSW model where the SPF data are utilized as conditioning information relative to the RMSE of the benchmark GSW model. This means that for values below (above) unity, the SPF-based GSW model has a lower (higher) RMSE than the benchmark GSW model. A few findings are worth highlighting. First, the relative RMSEs are typically lower than unity for real GDP growth, inflation, and unemployment, suggesting that the SPF data are useful for improving the

forecasts of these variables. Second, the relative RMSEs for the two cases of the GSW model that are subject to the news interpretation (dashed and dotted lines in Fig. 2) are similar, and may reflect the fact that, at best, the information in the 2-year-ahead SPF leads to a marginal improvement in the forecasting performance of the GSW model. Third, the GSW model subject to the noise interpretation (dash-dotted lines) appears to worsen the real wage forecasts and marginally improve the short-run employment forecasts. For the news cases, the opposite result is obtained. Fourth, all models using SPF data seem to worsen the interest rate forecasts. Fifth, consumption and investment forecasts may be improved marginally, with a tendency for the news models to fare better than the noise model.

To examine the findings in more detail, we turn to Table 4, where percentile values from the approximating distribution of the modified Diebold–Mariano test statistic are displayed; see Harvey, Leybourne, and Newbold (1997) for computational details. We have opted to follow the suggestion of Harvey et al. and compare the modified statistic to the Student's t -distribution with $N_h - 1$ degrees of freedom, where N_h is the number of h -step-ahead forecast errors, rather than to its asymptotic normal distribution. A low percentile value indicates that the SPF-based forecasts (noise or news) are better, while a high percentile value favors the forecasts of the benchmark model. A value below or equal to 5% is given in bold in the table, while a value above or equal to 95% is given in italics. With respect to the five observations listed in the previous paragraph, the results in Table 4 support the finding that, from a RMSE perspective, the inflation forecasts are improved for both the news models and the noise model when the SPF is taken into account. Furthermore, the news models seem to help improve the short-run real GDP growth forecasts, while the evidence on unemployment suggests that the forecasts are neither improved nor worsened when using the SPF data.

The results in Table 4 also support the third point above, namely that the SPF information is useful for improving real wage forecasts for the news model, but that this data set worsens the real wage forecasts for the noise model. Furthermore, the interest rate forecasts are generally worsened when the SPF data are included, especially for the

¹⁵ An alternative to the noise and news implementations of the SPF data would be to employ the methodology developed by Monti (2010). While her approach has several attractive properties, we have opted not to use it, for various reasons. First, the monthly vintage that we have selected to represent the quarter (third month of each quarter) is not consistent with the condition that judgmental forecasts are based on an information set which comprises the information available in the RTDB vintage. For the euro area RTDB, this condition could only be satisfied if the first month of each quarter were used, while our decision to use the third month is based on using a vintage that basically includes information which is available when the ECB/Eurosystem staff macroeconomic projections exercises are conducted. Second, her approach involves replacing the unknown population moments of the Kalman filter with sample moments calculated using the judgmental forecasts. The SPF data set has a very short historical sample for the first vintages in the forecasting study, namely eight observations on the judgmental forecasts for the 2001Q1 vintage. It therefore seems unlikely that the unknown population moments required by her approach can be estimated meaningfully for this vintage and those that immediately follow.

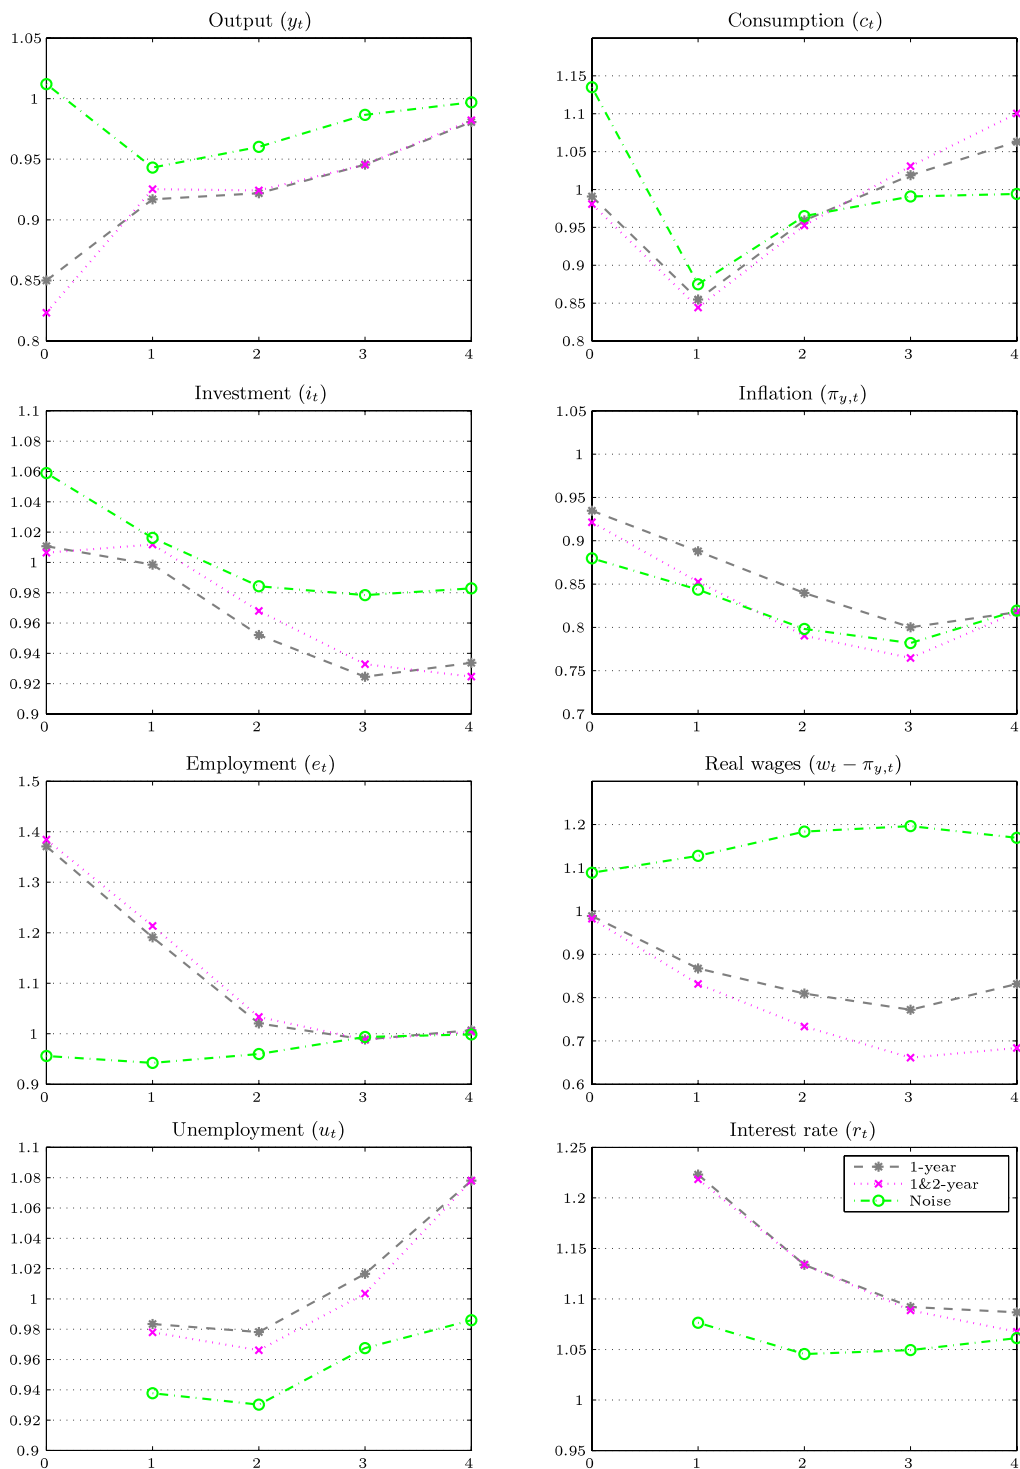


Fig. 2. Relative RMSEs for DSGE models when conditioning on SPF data, compared with RMSEs for the DSGE model without the SPF. The calculations are based on the RTDB vintages 2001Q1–2010Q4.

shorter forecast horizons. Overall, it would appear from the modified Diebold–Mariano tests that the SPF data are useful for improving the benchmark GSW model point forecasts under the news interpretation, while the evidence for

the noise case is less convincing. It should be kept in mind that each test is based on $N_h = 36 - h$ observations for the h -step-ahead forecasts, and should therefore be interpreted with caution, as the reference distribution need not

Table 5

Modified Diebold–Mariano tests of structural and non-structural models versus the BVAR model (RTDB vintages 2001Q1–2010Q4).

Variable	Benchmark GSW				Noise				News 1-year				News 1 and 2-year			
	1	2	3	4	1	2	3	4	1	2	3	4	1	2	3	4
Output	0.23	0.36	0.49	0.65	0.13	0.17	0.39	0.73	0.04	0.09	0.18	0.67	0.04	0.09	0.18	0.72
Consumption	<i>1.00</i>	<i>1.00</i>	<i>0.99</i>	<i>0.99</i>	<i>0.98</i>	<i>1.00</i>	<i>1.00</i>	<i>1.00</i>	<i>0.99</i>	<i>1.00</i>	<i>0.99</i>	<i>0.99</i>	<i>0.99</i>	<i>1.00</i>	<i>0.99</i>	<i>0.98</i>
Investment	0.21	0.27	0.32	0.39	0.21	0.17	0.15	0.16	0.19	0.20	0.21	0.27	0.22	0.22	0.22	0.26
Inflation	0.68	0.26	0.12	0.10	0.03	0.05	0.02	0.04	0.07	0.06	0.02	0.02	0.04	0.04	0.02	0.02
Employment	0.89	0.88	0.89	0.94	0.86	0.86	0.85	0.86	0.93	0.88	0.86	0.86	0.94	0.88	0.86	0.85
Real wages	<i>0.98</i>	<i>1.00</i>	<i>1.00</i>	<i>1.00</i>	<i>0.99</i>	<i>1.00</i>	<i>1.00</i>	<i>1.00</i>	0.76	0.94	<i>0.98</i>	<i>1.00</i>	0.59	0.71	0.65	0.80
Unemployment	0.34	0.48	0.60	0.59	0.06	0.20	0.52	0.59	0.20	0.33	0.87	0.91	0.21	0.29	0.81	0.91
Interest rate	0.10	0.15	0.21	0.16	0.24	0.20	0.27	0.20	0.64	0.36	0.33	0.21	0.63	0.36	0.33	0.20
	AR(1)				RW				Mean							
Output	0.89	0.87	0.83	0.81	0.92	0.92	0.89	0.89	0.90	0.85	0.81	0.76				
Consumption	<i>0.97</i>	<i>0.84</i>	<i>0.72</i>	<i>0.75</i>	<i>0.94</i>	<i>0.77</i>	<i>0.49</i>	<i>0.58</i>	<i>0.96</i>	<i>0.90</i>	<i>0.85</i>	<i>0.83</i>				
Investment	0.81	0.77	0.75	0.73	0.83	0.80	0.79	0.79	0.91	0.84	0.78	0.72				
Inflation	0.20	0.21	0.08	0.06	0.24	0.24	0.11	0.08	0.28	0.10	0.04	0.04				
Employment	0.94	0.92	0.90	0.90	0.95	0.93	0.92	0.92	0.93	0.88	0.84	0.81				
Real wages	0.01	0.05	0.09	0.16	0.03	0.09	0.12	0.19	0.01	0.03	0.08	0.15				
Unemployment	0.82	0.84	0.86	0.92	0.81	0.84	0.87	0.93	<i>1.00</i>	<i>0.99</i>	<i>0.98</i>	<i>0.97</i>				
Interest rate	0.65	0.39	0.32	0.16	0.61	0.38	0.32	0.17	<i>1.00</i>	<i>1.00</i>	<i>0.97</i>	<i>0.77</i>				

Notes: The modified Diebold–Mariano test has been calculated as per Harvey et al. (1997, equation (9)) for the squared forecast errors of the model displayed in the header of each column group versus the squared forecast errors of the BVAR model. Percentile values taken from the Student's t -distribution with $N_h - 1$ degrees of freedom are shown above, with N_h being the number of h -step-ahead forecast errors, $N_h = 36 - h$. Smaller percentile values favor the model displayed in the header, and larger percentile values favor the BVAR model. Numbers in bold refer to percentile values which are less than or equal to 5%, and numbers in italics to values greater than or equal to 95%.

provide a good approximation of the unknown small sample distribution of the test statistic.

Turning to the second question about how well the GSW model can compete with the non-structural models, we first consider the RMSEs reported in Fig. 3, which are all relative to the RMSEs of the BVAR model. In other words, values above unity favor the BVAR and values below unity the specified model. Concerning the four GSW model based cases (denoted DSGE, noise, 1-year and 1&2-year in the graph), it can be seen that they generally have lower RMSEs for inflation than the BVAR, especially when the SPF forecasts are included, and for the interest rate. On the other hand, the BVAR has lower RMSEs for consumption, employment and real wages, although the news model based on the 1- and 2-year-ahead SPF comes close for wages. Compared with the non-structural models, the BVAR has lower RMSEs for output, employment, and unemployment, while the non-structural models have better point forecasts for inflation and real wages.

Table 5 displays the percentile values of the modified Diebold–Mariano statistics for the four GSW model based cases with and without the SPF data relative to the BVAR in the upper half, and those for the three univariate non-structural models versus the BVAR in the lower half. Our findings based on the RMSEs are confirmed to some extent, especially for the comparisons between the GSW model based cases and the BVAR. The inflation forecasts of the GSW model cases seem to outperform the BVAR when the SPF conditioning information is utilized, while consumption in particular, as well as real wages, are generally forecasted better using the BVAR. As was noted above, the news case with 1- and 2-year-ahead SPF data seems to improve the real wage forecast sufficiently relative to the benchmark GSW and the news case using only the 1-year-ahead SPF data to match the point forecasts of the BVAR

from a MSE perspective. An inspection of the forecast errors reveals that the GSW model systematically overpredicts real wage growth and consumption. A similar result was found by Christoffel, Coenen, and Warne (2011), who evaluated the forecast performance of the NAWM for the euro area; see also Warne, Coenen, and Christoffel (2013). The New Keynesian model, which assumes a constant steady-state labor share and consumption to output ratio, has a difficult time explaining the falling labor share and the rising consumption to GDP ratio over this period. The non-structural models (except for the sample mean) do better in this respect, especially for the one- and two-step-ahead point forecasts. The noise versus news interpretation does matter for the predictive performance regarding wage growth. In the news model, the higher inflation HICP forecasts are rationalised by higher expected markup shocks, which at the same time tend to reduce the expected wage growth, thereby alleviating part of the upward bias of the benchmark DSGE model. On the other hand, in the noise model, the overprediction of real wage growth is magnified.

The graphs in Fig. 4 plot the log-determinant and the trace statistic of the MSE matrix for the four GSW model based cases with and without the SPF data, the AR(1) model, the sample mean, and the random walk model relative to the values obtained for the BVAR model. Concerning the log-determinant statistic, it should be noted that it is negative for all models, and we have therefore opted to compute the relative statistic as minus unity times the ratio of log-determinants. Hence, a low value suggests better point forecasts from a multivariate perspective, with minus one being the log-determinant of the BVAR model. Similarly, low values of the trace statistic are also an indicator of better multivariate point forecasts, with values below unity indicating that the trace of the MSE matrix of the given model is lower than the corresponding trace statistic of the BVAR model.

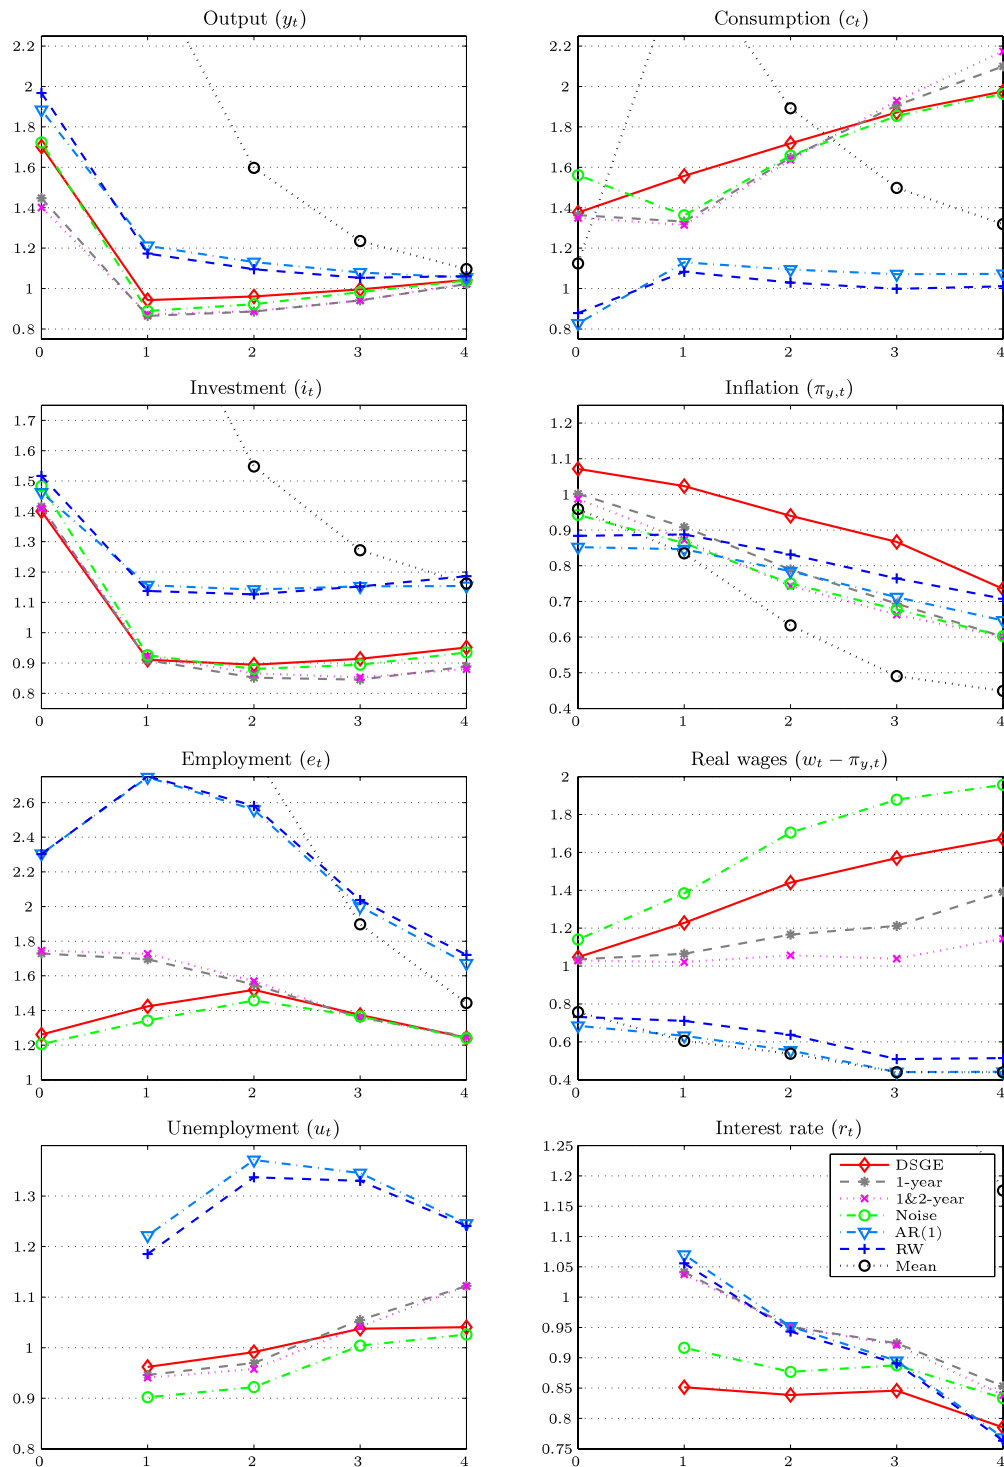


Fig. 3. Relative RMSEs of structural and non-structural models compared with RMSEs for the BVAR model. The calculations are based on the RTDB vintages 2001Q1–2010Q4.

From the trace statistics on the right hand side in Fig. 4, it would seem that the point forecasts of the GSW model based cases under the news interpretation do (marginally) better than all other models for all horizons, while the point forecasts of the benchmark GSW model and noise

broadly match those of the BVAR. The three simple non-structural models all have trace statistics greater than the BVAR, especially the sample mean for the shorter forecast horizons. Turning to the log-determinant on the left hand side of Fig. 4, the picture is far more complex. For one-

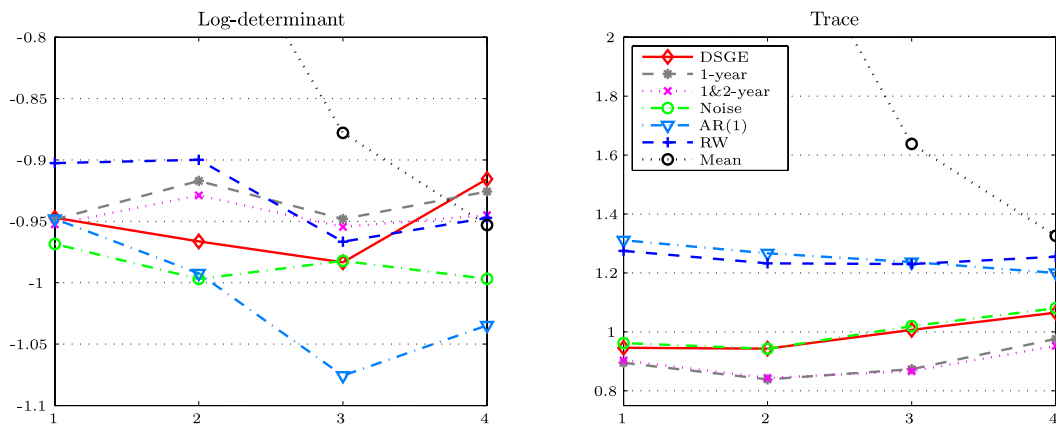


Fig. 4. Multivariate MSE statistics for the RTDB vintages 2001Q1–2010Q4.

step-ahead forecasts the BVAR does better, while the AR(1) and the noise case of the GSW model have statistics close to those of the BVAR for two-step-ahead forecasts. For the three- and four-step-ahead forecasts, the AR(1) model seems to perform best according to this metric.

To summarize the findings on the relative forecasting performances of the models, it appears that the SPF data are useful for improving the point forecasts of the GSW model, particularly for the cases which are subject to the news interpretation. The main cost concerns the forecasts of the short-term nominal interest rate at the one- to three-step-ahead horizons. Once the point forecasts of the GSW model are compared with those of non-structural models, the results are mixed. Compared with a BVAR model which does not take the SPF data into account, the news models improve the inflation forecasts and the one-step-ahead real GDP growth forecasts. However, these results seem to be driven mainly by the utilization of the SPF, rather than by having a structural model. Hence, a BVAR which takes the SPF data on board when forecasting could easily do at least as well as the news models when forecasting these variables. At the same time, the real wage point forecasts of the BVAR can be improved on, particularly when comparing them to the forecasts from the three univariate non-structural models.

The final exercise we shall conduct concerns how good the forecasts of these models are from an absolute perspective. To study this issue we follow Edge and Gürkaynak (2010) and Mincer and Zarnowitz (1969), for example, and conduct so-called Mincer–Zarnowitz regressions. Let x_t be the (actual) value of the forecasted variable at time t , while the h -step-ahead forecast of this variable is denoted by $x_{t|t-h}^{(m)}$ for model m . The regression equation is now

$$x_t = \alpha_h^{(m)} + \beta_h^{(m)} x_{t|t-h}^{(m)} + \varepsilon_{h,t}^{(m)}, \quad t = 1, \dots, N_h,$$

where $\varepsilon_{h,t}^{(m)}$ is a mean zero error term. If the forecasts of model m are efficient, then the intercept is zero while the slope is unity, and the variance of the error term is small compared with the variance of x_t (high R^2). An intercept which is different from zero indicates that, on average, the forecasts have been biased (relative to the selected actual);

and if the slope coefficient is greater (less) than unity, then the forecasts have consistently underpredicted (overpredicted) the variable. A low R^2 value indicates that little of the variation in x_t is captured by the variation of the forecast.

Table 6 summarizes the evidence from the forecast accuracy regressions, limiting the presentation to three models (the benchmark GSW model, the news model with 1-year-ahead SPF data, and the BVAR model), and four variables (output, inflation, unemployment, and the interest rate).¹⁶ The results suggest that forecasts of inflation, unemployment, and the interest rate have been poor by all methods. For output, however, the regression outcomes are broadly consistent with efficient forecasts, especially for the BVAR and the longer forecast horizons.

6. Conclusion

In this paper we have evaluated the real-time forecasting performance of the New Keynesian model of Galí et al. (2012), estimated on euro area data. First, we find that the benchmark GSW model forecasts tend to be improved by adding one- to two-year-ahead professional forecasts of real GDP growth, inflation, and the unemployment rate to the conditioning data without otherwise changing the DSGE model. This is consistent with the news interpretation, where it is assumed that the forecasts reveal the presence of expected future structural shocks in line with those estimated in the past. The consequence of utilizing the professional forecasts within the context of the GSW model is that inflation and real wage forecasts improve considerably, as do one- and two-quarter-ahead real GDP forecasts.

¹⁶ The standard errors in parentheses have been computed using the Newey and West (1987) method under the assumption that the error term follows an MA(h) process, except for unemployment and the interest rate, where it is assumed to follow an MA($h - 1$) process. This is due to the fact that, excluding the unemployment rate and the interest rate, all other variables need to be nowcasted for some vintages. There are three such vintages for output, seven vintages for consumption and investment, 14 vintages for inflation, and all 36 vintages for employment and real wages. We have also made use of this information when computing the modified Diebold–Mariano test.

Table 6

Forecast accuracy regressions for selected variables: the GSW model with and without SPF and BVAR (RTDB vintages 2001Q1–2010Q4).

Forecast	Benchmark GSW				News 1-year				BVAR			
	1	2	3	4	1	2	3	4	1	2	3	4
Output slope	1.17 (0.09)	1.26 (0.24)	1.25 (0.54)	0.67 (0.71)	1.11 (0.08)	1.30 (0.20)	1.90 (0.50)	1.57 (1.19)	1.05 (0.10)	1.07 (0.24)	1.59 (0.67)	1.26 (1.56)
Intercept	−0.35 (0.18)	−0.83 (0.45)	−1.26 (1.06)	−0.42 (1.52)	−0.27 (0.17)	−0.88 (0.38)	−2.36 (0.94)	−2.26 (2.47)	−0.23 (0.21)	−0.53 (0.47)	−1.68 (1.19)	−1.28 (2.77)
Adj. R^2	0.08	0.30	0.56	0.77	0.07	0.23	0.39	0.71	0.11	0.36	0.57	0.79
Inflation slope	0.54 (0.16)	0.43 (0.24)	−0.27 (0.35)	−0.70 (0.51)	0.64 (0.17)	0.56 (0.27)	0.01 (0.40)	−0.60 (0.57)	0.48 (0.14)	0.24 (0.34)	−0.11 (0.86)	−0.13 (1.27)
Intercept	1.07 (0.28)	1.31 (0.41)	2.45 (0.57)	3.15 (0.83)	0.87 (0.31)	1.04 (0.49)	2.01 (0.72)	3.13 (1.06)	1.15 (0.28)	1.59 (0.36)	2.22 (0.33)	2.26 (0.31)
Adj. R^2	−0.01	0.00	0.01	0.00	−0.01	−0.00	0.01	0.00	−0.00	0.00	0.01	0.01
Unemployment slope	0.84 (0.06)	0.83 (0.13)	0.74 (0.23)	0.53 (0.34)	0.83 (0.06)	0.78 (0.10)	0.68 (0.17)	0.50 (0.24)	0.79 (0.05)	0.75 (0.09)	0.69 (0.15)	0.59 (0.21)
Intercept	1.31 (0.53)	1.40 (1.07)	2.19 (1.94)	3.95 (2.83)	1.44 (0.49)	1.90 (0.85)	2.78 (1.41)	4.26 (1.97)	1.77 (0.45)	2.20 (0.75)	2.66 (1.21)	3.56 (1.70)
Adj. R^2	−0.03	−0.03	−0.03	−0.03	−0.03	−0.03	−0.03	−0.03				
Interest rate slope	0.93 (0.06)	0.80 (0.15)	0.56 (0.29)	0.24 (0.38)	0.98 (0.08)	0.80 (0.19)	0.50 (0.30)	0.22 (0.37)	0.81 (0.05)	0.64 (0.09)	0.51 (0.17)	0.28 (0.21)
Intercept	0.13 (0.19)	0.40 (0.53)	1.07 (1.04)	2.13 (1.40)	−0.10 (0.26)	0.32 (0.67)	1.19 (1.12)	2.13 (1.44)	0.42 (0.16)	0.80 (0.35)	1.04 (0.67)	1.86 (0.90)
Adj. R^2	−0.02	0.02	0.06	0.08	−0.01	0.03	0.06	0.08	−0.02	0.00	0.04	0.07

The cost appears to be that the short-term nominal interest rate forecasts deteriorate. The evidence under the noise interpretation, where the mean professional forecasts are assumed to be noisy indicators of the rational expectations forecasts implied by the DSGE model, is less convincing. Although inflation forecasts seem to be improved under the noise model, real wage forecasts seem to deteriorate, as do the short-term nominal interest rate forecasts.

Second, a BVAR model is also able to improve the benchmark GSW model forecasts, particularly for consumption and real wages where the DSGE model systematically overpredicts real wage growth and consumption. These variables are also predicted poorly under the noise and news interpretations of the GSW model relative to the BVAR. At the same time, the inflation forecasts from these models are typically improved considerably relative to the BVAR.

Third, the point forecasts of the variables are generally not efficient when viewed through the lens of Mincer–Zarnowitz regressions. One exception is output, where the regression evidence for most of the models is consistent with a zero intercept, a unit slope coefficient, and a relatively high adjusted R^2 value. It should be kept in mind that the forecast performance study covers a sample with 36 real-time vintages from the ECB's Statistical Data Warehouse. In view of this rather small sample, the evidence needs to be interpreted with caution. Furthermore, we only study point forecasts and do not take other moments of the predictive distributions into account.

Acknowledgments

We are grateful to two anonymous referees and the handling editor Graham Elliott for useful comments. We have also benefitted from comments by Jonathan Wright and Francesca Monti, as well as participants at the American Economic Association Annual Meeting in Chicago, January 2012, the 7th ECB Workshop on Forecasting Techniques, May 2012, and the SPF Workshop at the European Central Bank, November 2012. The views expressed are our own and should not be attributed to the European Central Bank, the National Bank of Belgium or the Eurosystem.

References

- Altissimo, F., Ehrmann, M., & Smets, F. (2006). *Inflation persistence and price-setting behaviour in the euro area: a summary of the IPN evidence*. ECB occasional paper series No. 46. Frankfurt am Main: European Central Bank.
- Bowles, C., Friz, R., Genre, V., Kenny, G., Meyler, A., & Rautanen, T. (2007). *The ECB Survey of Professional Forecasters (SPF): a review after eight years' experience*. ECB occasional paper series No. 59. Frankfurt am Main: European Central Bank.
- Červená, M., & Schneider, M. (2014). Short-term forecasting of GDP with a DSGE model augmented by monthly indicators. *International Journal of Forecasting*, 30, 498–516.
- Christoffel, K., Coenen, G., & Warne, A. (2011). Forecasting with DSGE models. In M. P. Clements & D. F. Hendry (Eds.), *The Oxford handbook of economic forecasting* (pp. 89–127). New York: Oxford University Press.
- Coenen, G., Levin, A., & Wieland, V. (2005). Data uncertainty and the role of money as information variable for monetary policy. *European Economic Review*, 49, 975–1006.

- Coenen, G., & Warne, A. (2013). *Risks to price stability, the zero lower bound and forward guidance: a real-time assessment*. ECB working paper series No. 1582. Frankfurt am Main: European Central Bank.
- Croushore, D. (2011). Frontiers of real-time data analysis. *Journal of Economic Literature*, 49, 72–110.
- Croushore, D., & Stark, T. (2001). A real-time data set for macro-econometricians. *Journal of Econometrics*, 105, 111–130.
- Del Negro, M., & Eusepi, S. (2011). Fitting observed inflation expectations. *Journal of Economic Dynamics and Control*, 35, 2105–2131.
- Del Negro, M., & Schorfheide, F. (2013). DSGE model-based forecasting. In G. Elliott & A. Timmermann (Eds.), *Handbook of economic forecasting*, Vol. 2 (pp. 57–140). Amsterdam: North-Holland.
- Diebold, F. X., & Rudebusch, G. D. (1991). Forecasting output with the composite leading index: a real-time analysis. *Journal of the American Statistical Association*, 86, 603–610.
- Durbin, J., & Koopman, S. J. (2012). *Time series analysis by state space methods* (2nd ed.). Oxford: Oxford University Press.
- Edge, R. M., & Gürkaynak, R. S. (2010). How useful are estimated DSGE model forecasts for central bankers? *Brookings Papers on Economic Activity*, Fall, 209–259.
- Fagan, G., Henry, J., & Mestre, R. (2005). An area-wide model for the euro area. *Economic Modelling*, 22, 39–59.
- Faust, J., & Wright, J. H. (2013). Forecasting inflation. In G. Elliott & A. Timmermann (Eds.), *Handbook of economic forecasting*, Vol. 2 (pp. 2–56). Amsterdam: North-Holland.
- Galí, J., Smets, F., & Wouters, R. (2012). Unemployment in an estimated new Keynesian model. In D. Acemoglu & M. Woodford (Eds.), *NBER macroeconomics annual 2011* (pp. 329–360). University of Chicago Press.
- Garcia, J.A. (2003). *An introduction to the ECB's survey of professional forecasters*. ECB occasional paper series: vol. 8. Frankfurt am Main: European Central Bank.
- Genre, V., Kenny, G., Meyler, A., & Timmermann, A. (2013). Combining expert forecasts: can anything beat the simple average? *International Journal of Forecasting*, 29, 108–121.
- Giannone, D., Henry, J., Lalik, M., & Modugno, M. (2010). *An area-wide real-time database for the euro area*. ECB working paper series No. 1145.
- Giannone, D., Henry, J., Lalik, M., & Modugno, M. (2012). *An area-wide real-time database for the euro area*. *Review of Economics and Statistics*, 94, 1000–1013.
- Giannone, D., Monti, F., & Reichlin, L. (2010). Incorporating conjunctural analysis in structural models. In V. Wieland (Ed.), *The science and practice of monetary policy today* (pp. 41–57). Springer-Verlag.
- Greenwood, J., Hercowitz, Z., & Huffman, G. (1988). Investment, capacity utilization and the real business cycle. *American Economic Review*, 78, 402–417.
- Harvey, D., Leybourne, S., & Newbold, P. (1997). Testing the equality of prediction mean squared errors. *International Journal of Forecasting*, 13, 281–291.
- Mincer, J. A., & Zarnowitz, V. (1969). The evaluation of economic forecasts. In J. A. Mincer (Ed.), *Economic forecasts and expectations: analysis of forecasting behavior and performance* (pp. 3–46). NBER.
- Monti, F. (2010). Combining judgment and models. *Journal of Money, Credit and Banking*, 42, 1641–1662.
- Newey, W. K., & West, K. D. (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55, 703–708.
- Smets, F., & Wouters, R. (2003). An estimated stochastic dynamic general equilibrium model for the euro area. *Journal of the European Economic Association*, 1, 1123–1175.
- Smets, F., & Wouters, R. (2007). Shocks and frictions in US business cycles: a Bayesian DSGE approach. *American Economic Review*, 97, 586–606.
- Villani, M. (2009). Steady-state priors for vector autoregressions. *Journal of Applied Econometrics*, 24, 630–650.
- Waggoner, D. F., & Zha, T. (1999). Conditional forecasts in dynamic multivariate models. *Review of Economics and Statistics*, 81, 639–651.
- Warne, A. (2013). *YADA manual—computational details*. Manuscript, European Central Bank. Available with the YADA distribution.
- Warne, A., Coenen, G., & Christoffel, K. (2013). *Predictive likelihood comparisons with DSGE and DSGE-VAR models*. ECB working paper series No. 1536. Frankfurt am Main: European Central Bank.
- WDN (2009). *Wage dynamics in Europe: final report of the wage dynamics network (WDN)*. Mimeo http://www.ecb.europa.eu/home/pdf/wdn_finalreport_dec2009.pdf.

Frank Smets is Adviser to the President of the European Central Bank. He is professor of international economics at the Centre for Economic Studies at the KU Leuven and an honorary professor in the Duisenberg chair at the Faculty of Economics and Business of the University of Groningen. He is a Research Fellow of the CEPR in London and CESifo in Munich. He has written and published extensively on monetary, macroeconomic, financial and international issues mostly related to central banking in top academic journals.

Anders Warne is a senior economist at the Monetary Policy Research Division in the Directorate General Research of the European Central Bank. His main research interests are time series analysis for applied macroeconomics and Bayesian methods.

Rafael Wouters is advisor to the research Department of the National Bank of Belgium. He is teaching at the Centre for Economic Studies at the KU Leuven and at Ecares at ULB. He is a Research Fellow of CEPR and member of the EABCN. He has published on the estimation and the use of DSGE models for monetary policy analysis.